

System calls for using TCP

Client

socket – make socket

bind – assign address (optional)

connect – connect to listening socket

write – send data

read – receive data

Server

socket – make socket

bind – assign address

listen – listen for clients

accept – accept connection

read – receive data

write – send data

- Anything **red** might block, waiting for network
 - Obviously bad for applications that need concurrency

Non-blocking I/O

- **Use `fcntl` to set `O_NONBLOCK` flag on descriptor**
- **Non-blocking semantics of system calls:**
 - `read` immediately returns -1 with `errno` `EAGAIN` if no data
 - `write` may not write all data, or may return `EAGAIN`
 - `connect` may “fail” with `EINPROGRESS` (or may succeed, or may fail with real error like `ECONNREFUSED`)
 - `accept` may fail with `EAGAIN` if no pending connections

How to know when to read/write?

```
struct pollfd {
    int fd;           /* file descriptor */
    short events;      /* Events you are interested in */
    short revents;     /* Events that have happened (results) */
};

int poll(struct pollfd *fds, nfds_t nfd, int timeout);

/* Some possible events: */
#define POLLIN      0x0001 /* Can read fd without blocking */
#define POLLOUT     0x0004 /* Can write fd without blocking */
#define POLLERR     0x0008 /* Error on fd (only in revents) */
#define POLLHUP     0x0010 /* 'Hangup' has occurred on fd */
```

- **Note: BSD used select to achieve same thing**
 - Most OSes support both select and poll today

epoll

- **Newer Linux provides** `epoll`
- **Interface allows more efficient implementation**
 - Register interest with `epoll_ctl` syscall
 - Wait with `epoll_wait` syscall
 - Kernel doesn't have to re-scan `pollfd` array on each wait
- **New option bits reduce calls to** `epoll_ctl`
 - `EPOLLONESHOT` – only wait for event once
 - `EPOLLET` – “edge triggered” (as opposed to level triggered)
- **`epoll` is Linux specific**
 - But BSD has `kqueue/kevent` which is similar idea

epoll interface

```
typedef union epoll_data {  
    int fd;  
    /* ... */  
} epoll_data_t;
```

```
struct epoll_event {  
    __uint32_t events;      /* Epoll events */  
    epoll_data_t data;      /* User data variable */  
};
```

```
int epoll_create(int size);  
int epoll_ctl(int epfd, int op, int fd,  
              struct epoll_event *event);  
int epoll_wait(int epfd, struct epoll_event *events,  
               int maxevents, int timeout);
```

Asynchronous programming model

- **Many non-blocking file descriptors in one process**
 - Wait for pending I/O events on file many descriptors
 - Each event triggers some *callback* function
- **E.g., build “callback harness”:**

```
/* Register callback for when fd is readable or writable */  
void cb_add (int fd, int write, void (*fn)(void *), void *arg);
```

```
/* Unregister callback */  
void cb_free (int fd, int write);
```

```
/* Loop forever checking callbacks */  
void cb_check (void);
```

Simplified example

```
struct state {
    int fd;
    /* ... */
};

void doit (void) {
    struct state *st = malloc (sizeof (*st));
    st->fd = create_new_tcp_socket ();
    connect (st->fd, &someplace, sizeof (someplace));
    cb_add (st->fd, 1, doit_2, st);
}

static void doit_2 (void *_st) {
    struct state *st = _st;
    write (st->fd, "request\n", 8);
    cb_free (st->fd, 1);
    cb_add (st->fd, 0, doit_3, st);
}

static void doit_3 (void *_st) {
    struct state *st = _st;
    /* read more from st->fd until you get full response */
}
```

Syntactic sugar

- **Problem: Need state from one callback to next**
- **E.g., C++ can implement `wrap` that bundles a function with its arguments**

```
callback<void, int>::ref errwrite = wrap (write, 2);  
(*errwrite) ("hello", 5); // calls write (2, "hello", 5);
```

- **Possible to build large event-driven apps this way**
 - E.g., I have built large library to do this
 - Debugging features include recording where callbacks created to facilitate tracing
- **Google reportedly does similar things**

Intro to Threads

- **Threads: most popular abstraction for concurrency**
 - Lighter-weight abstraction than processes
 - All threads in one process have same memory, file desc., etc.
 - Allows one process to use multiple CPUs
- **Example: threaded web server:**
 - Service many clients simultaneously

```
for (;;) {  
    fd = accept_client ();  
    thread_create (service_client, &fd);  
}
```

How to share CPU amongst threads

- **Each thread has execution state:**
 - Stack, program counter, registers, condition codes, etc.
- **Switch the CPU amongst the threads**
 - Save away execution state of one, load up that of next
- **When to switch?**
 - Current thread can no longer use the CPU (waiting for I/O)
 - Current thread has had CPU for too long (preemption)
 - Scheduler maintains lists of runnable/running/waiting threads

Thread package API

- `tid create (void (*fn) (void *), void *arg);`
 - Create a new thread, run fn with arg
- `void exit ();`
 - Destroy current thread
- `void join (tid thread);`
 - Wait for thread thread to exit

Synchronization primitives

- `void lock (mutex_t m);`
`void unlock (mutex_t m);`
 - Only one thread acquires `m` at a time, others wait
 - **All global data must be protected by a mutex!**
- `void wait (mutex_t m, cond_t c);`
 - Atomically unlock `m` and sleep until `c` signaled
- `void signal (cond_t c);`
`void broadcast (cond_t c);`
 - Wake one/all users waiting on `c`

Example: Taking job from work queue

```
job *job_queue;
mutex_t job_mutex;
cond_t job_cond;
void workthread (void *) {
    job *j;
    for (;;) {
        lock (job_mutex);
        while (!(j = job_queue))
            wait (job_mutex, job_cond);
        job_queue = j->next;
        unlock (job_mutex);
        do (j);
    }
}
```

Example: Adding job to work queue

```
void addjob (job *j) {  
    lock (job_mutex);  
    j->next = job_queue;  
    job_queue = j;  
    signal (job_cond);  
    unlock (job_mutex);  
}
```

- **Atomic release/wait necessary in** `workthread`, **otherwise:**
 - `workthread` checks queue, releases lock
 - `addjob` adds job to queue, signals `job_mutex`
 - `workthread` waits for signal that was already delivered

Other thread package features

- Alerts – cause exception in a thread
- Trylock – don't block if can't acquire mutex
- Timedwait – timeout on condition variable
- Shared locks – concurrent read accesses to data
- Thread priorities – control scheduling policy
- Thread-specific global data

Implementing shared locks

```
struct sharedlk {
    int i; mutex_t m; cond_t c;
};

void AcquireExclusive (sharedlk *sl) {
    lock (sl->m);
    while (sl->i) { wait (sl->m, sl->c); }
    sl->i = -1;
    unlock (sl->m);
}

void AcquireShared (sharedlk *sl) {
    lock (sl->m);
    while (sl->i < 0) { wait (sl->m, sl->c); }
    sl->i++;
    unlock (sl->m);
}
```


shared locks (continued)

```
void ReleaseShared (sharedlk *sl) {  
    lock (sl->m);  
    if (!--sl->i) signal (sl->c);  
    unlock (sl->m);  
}  
  
void ReleaseExclusive (sharedlk *sl) {  
    lock (sl->m);  
    sl->i = 0;  
    broadcast (sl->c);  
    unlock (sl->m);  
}
```

- **Must deal with starvation**

Deadlock

- **Mutex ordering:**
 - A locks m1, B locks m2, A locks m2, B locks m1
 - How to avoid?
- **Similar deadlock with condition variables**
 - Suppose resource 1 managed by c_1 , resource 2 by c_2
 - A has 1, waits on c_2 , B has 2, waits on c_1
- **Mutex/condition variable deadlock:**
 - `lock (a); lock (b); while (!ready) wait (b, c);
unlock (b); unlock (a);`
 - `lock (a); lock (b); ready = true; signal (c);
unlock (b); unlock (a);`

Moral: Bad to hold locks when crossing abstraction barriers!

Data races

- **Example: modify global ++x without mutex**
 - Might compile to: load, add 1, store
 - Bad interleaving changes result: load, load, ...
- **Even single instructions can have races**
 - E.g., `addl $1, _x`
 - Not atomic on MP without lock prefix!
- **Even reads dangerous on some architectures**
- **But sometimes cheating buys efficiency**

```
if (!initialized) {  
    lock (m);  
    if (!initialized) { initialize (); initialized = 1; }  
    unlock (m);  
}
```

Implementing user-level threads

- Allocate a new stack for each thread create
- Keep a queue of runnable threads
- Replace networking system calls (read/write/etc.)
 - If operation would block, switch and run different thread
- Schedule periodic timer signal (setitimer)
 - Switch to another thread on timer signals (preemption)

Example

- **Per-thread state in thread control block structure**

```
typedef struct tcb {  
    unsigned long md_esp;           /* Stack pointer of thread */  
    char *t_stack;                 /* Bottom of thread's stack */  
    /* ... */  
};
```

- **Machine-dependent thread-switch function:**
 - void thread_md_switch (tcb *current, tcb *next);
- **Machine-dependent thread initialization function:**
 - void thread_md_init (tcb *t,
 void (*fn) (void *), void *arg);

i386 thread_md_switch

```
pushl %ebp; movl %esp,%ebp          # Save frame pointer
pushl %ebx; pushl %esi; pushl %edi  # Save callee-saved regs

movl 8(%ebp),%edx                   # %edx = thread_current
movl 12(%ebp),%eax                  # %eax = thread_next
movl %esp, (%edx)                   # %edx->md_esp = %esp
movl (%eax), %esp                   # %esp = %eax->md_esp

popl %edi; popl %esi; popl %ebx     # Restore callee saved regs
popl %ebp                          # Restore frame pointer
ret                                 # Resume execution
```

i386 thread_md_init

```
void thread_md_init (tcb *t, void (*fn) (void *), void *arg) {
    u_long *sp = (u_long *) (t->t_stack + thread_stack_size);

    /* Set up a callframe to thread_begin */
    *--sp = (u_long) arg;      *--sp = (u_long) fn;
    *--sp = (u_long) t;        *--sp = 0; /* No return address */

    /* Now set up saved registers for switch.S */
    *--sp = (u_long) thread_begin; /* return address */
    *--sp = 0; /* ebp */          *--sp = 0; /* ebx */
    *--sp = 0; /* esi */          *--sp = 0; /* edi */

    t->t_md.md_esp = (mdreg_t) sp;
}
```

- **Swich will call thread_begin (fn, arg);**

Implementing kernel level threads

- **Start with process abstraction in kernel**
- **Strip out unnecessary features**
 - Same address space
 - Same file table
 - (Plan9's `rfork` actually allows individual control)
- **Faster than a process, but still very heavy weight**